

CSAI 301 - Project Phase 2

Connect 4 Game

Ahmed Wael 202201415

Game Description: Connect Four

Game Overview

Connect Four is a classic two-player deterministic game played on a vertical 6×7 grid. Players alternate dropping colored discs into columns, with gravity pulling each disc to the lowest available position. The objective is to form a horizontal, vertical, or diagonal line of four consecutive discs before the opponent does.

Why Connect Four for Adversarial Search?

- **Deterministic:** No randomness; perfect information available to both players
- **Two-player zero-sum:** One player's gain is the other's loss
- **Finite state space:** Game always terminates within 42 moves
- **Strategic complexity:** Requires lookahead planning and threat analysis
- **Computational tractability:** Suitable for depth-limited search with pruning

State Representation

Board Structure:

	A	B	C	D	E	F	G
1	●	●	●	●	●	●	●
2	●	●	●	●	●	●	●
3	●	●	●	●	●	●	●
4	●	●	●	●	●	●	●
5	●	●	●	●	●	●	●
6	●	●	●	●	●	●	●
7	●	●	●	●	●	●	●

Implementation:

- **Data Structure:** 6×7 NumPy array
- **Cell Values:**
 - 0 = Empty cell
 - 1 = Player 1 (Maximizing player/AI)
 - 2 = Player 2 (Minimizing player/Opponent)

Initial State

The game begins with:

- An empty 6×7 board (all cells = 0)
- Player 1 designated to move first
- No pieces on the board
- Empty move history

Actions and Moves

Valid Actions: At any states, the set of legal actions consists of all columns that are not filled

Properties:

- Minimum actions: 0 (board full - terminal state)
- Maximum actions: 7 (empty board)

Move Execution: When a player selects column c:

1. Find lowest empty row in column c
2. Place player's disc at position (row, c)
3. Update game state
4. Switch player

End game States and Utility

End game Conditions:

1. Win Condition: Four consecutive discs of the same player in any direction:

- **Horizontal, Vertical, Diagonal**

2. Draw Condition:

- Board full and No player has achieved win condition

3. Loss Condition:

- Opponent achieves four-in-a-row

Game Tree Complexity

Branching Factor Analysis:

- **Initial state:** $b = 7$ (all columns available)
- **Average state:** $b \approx 4-5$ (some columns filled)
- **Late game:** $b \approx 2-3$ (most columns near full)
- **Effective branching factor:** $\bar{b} \approx 4$

Depth Analysis:

- **Maximum depth:** $d = 42$ (all cells filled)
- **Average game length:** $\bar{d} \approx 35-40$ moves

Why Appropriate for Minimax & Alpha-Beta:

1. **Finite horizon:** Game always terminates
2. **Perfect information:** Both players see entire board
3. **Zero-sum property:** Enables minimax principle
4. **Reasonable branching:** $b \approx 4$ is manageable with pruning
5. **Evaluation feasible:** Heuristics can approximate position strength

Minimax Algorithm

Description

Minimax is a recursive algorithm that explores the game tree to find the optimal move for the current player. It operates on the principle that:

- **MAX player** (Player 1) tries to maximize the evaluation score
- **MIN player** (Player 2) tries to minimize the evaluation score

Core Concept: MAX assumes MIN will play optimally, and vice versa. The algorithm recursively evaluates all possible move sequences up to a depth limit, then selects the move leading to the best guaranteed outcome.

Implementation Details

Key Features:

1. **Depth Limiting:** Search stops at predetermined depth
2. **Terminal Detection:** Checks for win/loss/draw conditions
3. **Move Ordering:** Prioritizes center columns for better efficiency
4. **Performance Tracking:** Counts nodes explored and time taken

Advantages and Limitations

Advantages:

- ✓ Guaranteed to find optimal move
- ✓ Simple to understand and implement
- ✓ Always correct with sufficient depth

Limitations:

- ✗ Explores all nodes (no pruning)
- ✗ Exponential time complexity: $O(b^d)$
- ✗ Slow for deeper searches
- ✗ Examines obviously bad moves

Alpha-Beta Pruning Algorithm

Description

Alpha-Beta Pruning is an optimization of Minimax that eliminates branches in the game tree that cannot possibly influence the final decision. It maintains two values:

- **Alpha (α):** Best value MAX can guarantee so far
- **Beta (β):** Best value MIN can guarantee so far

Key Insight: If MIN finds a move worse (for MAX) than what MAX can already achieve elsewhere, there's no need to explore further in that branch.

Pruning Conditions

Beta Cutoff (in MAX node):

```
if β ≤ α:  
    prune remaining branches  
    (MIN already has better option elsewhere)
```

Alpha Cutoff (in MIN node):

```
if β ≤ α:  
    prune remaining branches  
    (MAX already has better option elsewhere)
```

Implementation Features

Optimizations:

1. **Move Ordering:** Center columns examined first (better pruning)
2. **Pruning Counter:** Tracks number of cutoffs
3. **Metrics Tracking:** Performance comparison with Minimax
4. **Same Interface:** Drop-in replacement for Minimax

3.5 Complexity Analysis

Time Complexity:

- **Best case:** $O(b^{(d/2)})$ - perfect move ordering

- **Average case:** $O(b^{(3d/4)})$ - random ordering
- **Worst case:** $O(b^d)$ - worst possible ordering

Space Complexity:

- $O(d)$ - recursive call stack

Practical Impact:

- Depth 6 Minimax: $\sim 4^6 = 4,096$ nodes
- Depth 6 Alpha-Beta: $\sim 4^3 = 64$ nodes (best case)
- **Effective doubling of searchable depth**

Evaluation Function

Purpose

Since searching to terminal states is infeasible (depth 40+), we need a heuristic evaluation function to estimate the "goodness" of non-terminal positions.

Design Goals:

1. Fast computation (called millions of times)
2. Accurate position assessment
3. Captures strategic features
4. Differentiates winning vs. losing positions

Evaluation Components

1. Window Counting

A "window" is any sequence of 4 consecutive positions (horizontal, vertical, or diagonal).

Scoring:

Four-in-a-row (win):	+1000
Three-with-one-empty:	+10
Two-with-two-empty:	+2
Blocked (mixed pieces):	0

Center Column Control

The center column (column 3) is strategically valuable as it participates in more potential winning lines.

Scoring:

```
Each piece in center column: +3 points
```

Rationale:

- Center creates more four-in-a-row possibilities
- Controls board geography
- Empirically strong in Connect Four strategy

2. Opponent Evaluation

The evaluation considers both players:

```
Score = (Player_Features) - (Opponent_Features)
```

This creates a relative advantage metric rather than absolute position strength.

Why This Function Works

1. Captures Immediate Threats:

- Three-in-a-row scores +10, making it valuable
- AI will prioritize creating threats or blocking opponent threats

2. Strategic Positioning:

- Two-in-a-row gets +2
- Center control bonus promotes strong positions

3. Balanced Evaluation:

- Considers both offensive and defensive features
- Relative scoring prevents bias

4. Computational Efficiency:

- Linear scan through board: $O(\text{rows} \times \text{cols})$
- No complex pattern matching

Limitations and Future Improvements

Current Limitations:

- Doesn't detect complex forced-win sequences beyond depth
- Equal weight to all three-in-a-row (some are stronger)
- No consideration of "traps" (multi-threat positions)

Potential Improvements:

- **Threat detection:** Immediate win/block recognition
- **Pattern library:** Known strong/weak configurations
- **Dynamic weights:** Adjust based on game phase

Performance Comparison

```
=====
Test 1: Early game (few moves)
=====
```

RESULTS

```
=====
```

Criterion	Minimax	Alpha-Beta
Best Move	Column 3	Column 3
Time Taken (s)	16.9490	0.5710
Nodes Explored	19608	778
Max Depth Reached	5	5
Best Score	12	12
Branches Pruned	N/A	131

```
=====
```

EFFICIENCY ANALYSIS

```
=====
```

Time Improvement: 96.63% faster
 Nodes Reduction: 96.03% fewer nodes
 Pruning Efficiency: 131 branches pruned
 Same Best Move: True

```
=====
```

Test 2: Mid game

=====

RESULTS

=====

Criterion	Minimax	Alpha-Beta
Best Move	Column 5	Column 5
Time Taken (s)	14.1387	4.4714
Nodes Explored	15699	5374
Max Depth Reached	5	5
Best Score	-8	-8
Branches Pruned	N/A	662

=====

EFFICIENCY ANALYSIS

=====

Time Improvement: 68.37% faster

Nodes Reduction: 65.77% fewer nodes

Pruning Efficiency: 662 branches pruned

Same Best Move: True

=====

Test 3: Complex position

=====

RESULTS

=====

Criterion	Minimax	Alpha-Beta
Best Move	Column 2	Column 2
Time Taken (s)	112.5486	11.6636
Nodes Explored	130607	14210
Max Depth Reached	6	6
Best Score	2	2
Branches Pruned	N/A	2458

=====

EFFICIENCY ANALYSIS for Alpha-Beta

=====

Time Improvement: 89.64% faster


Nodes Reduction: 89.12% fewer nodes

Pruning Efficiency: 2458 branches pruned

Same Best Move: True

Visual Results

Minimax vs Alpha-Beta



Performance Comparison Visualization

Minimax vs Alpha-Beta: Comprehensive Performance Analysis

